

21AIE311 – Reinforcement Learning

J Viswaksena - AM.EN.U4AIE21035

1.

```
import pygame
import sys
import numpy as np
import time
from collections import deque

# Initialize Pygame
pygame.init()

# Constants
WIDTH, HEIGHT = 500, 500 # Adjusted width and height
ROWS, COLS = 10, 10
CELL_SIZE = WIDTH // COLS # Adjusted cell size
GRID_COLOR = (0, 0, 0)
BACKGROUND_COLOR = (255, 255, 255)
SPECIAL_CELL_COLOR_YELLOW = (255, 255, 0)
SPECIAL_CELL_COLOR_RED = (255, 0, 0)
SPECIAL_CELL_COLOR_GREEN = (0, 255, 0)
SPECIAL_CELLS_YELLOW = [(0, 8), (0, 9), (1, 8), (1, 9)]
SPECIAL_CELLS_RED = [(2, 2), (2, 3), (2, 4), (2, 5), (2, 6), (3, 2), (4, 2), (5, 2), (6, 2), (7, 2),
                      (8, 2), (3, 6), (4, 6), (5, 6), (6, 6), (7, 6), (8, 6), (8, 5), (8, 3)]
CELL_VALUES = {}

# Set values for special cells
for cell in SPECIAL_CELLS_YELLOW:
    CELL_VALUES[cell] = 20 # Reward for reaching the goal

for cell in SPECIAL_CELLS_RED:
    CELL_VALUES[cell] = -1 # Penalty for hitting red cells

# Set value for all remaining cells
for row in range(ROWS):
    for col in range(COLS):
        if (row, col) not in CELL_VALUES:
            CELL_VALUES[(row, col)] = 1 # Default value for white cells

# Map actions to movements
ACTIONS = {"up": (-1, 0), "down": (1, 0), "left": (0, -1), "right": (0, 1)}

# Set up the display
```

```

win = pygame.display.set_mode((WIDTH, HEIGHT))
pygame.display.set_caption("Reinforcement Learning on Grid")

def draw_grid():
    for x in range(0, WIDTH, CELL_SIZE):
        pygame.draw.line(win, GRID_COLOR, (x, 0), (x, HEIGHT))
    for y in range(0, HEIGHT, CELL_SIZE):
        pygame.draw.line(win, GRID_COLOR, (0, y), (WIDTH, y))

def draw_cell(row, col):
    cell_rect = pygame.Rect(col * CELL_SIZE, row * CELL_SIZE, CELL_SIZE, CELL_SIZE)
    if (row, col) in SPECIAL_CELLS_YELLOW:
        pygame.draw.rect(win, SPECIAL_CELL_COLOR_YELLOW, cell_rect)
    elif (row, col) in SPECIAL_CELLS_RED:
        pygame.draw.rect(win, SPECIAL_CELL_COLOR_RED, cell_rect)
    else:
        pygame.draw.rect(win, BACKGROUND_COLOR, cell_rect)

def draw_agent(agent_position):
    agent_rect = pygame.Rect(agent_position[1] * CELL_SIZE, agent_position[0] * CELL_SIZE,
CELL_SIZE, CELL_SIZE)
    pygame.draw.rect(win, SPECIAL_CELL_COLOR_GREEN, agent_rect)

def main():
    agent_position = (4, 4) # Initial position of the agent
    visited_cells = set() # Keep track of visited cells
    queue = deque([agent_position]) # Queue for breadth-first search
    parent = {agent_position: None} # Parent dictionary for storing the shortest path
    reached_yellow_cell = False # Flag to indicate if the agent reached a yellow cell

    while queue:
        current_position = queue.popleft()
        visited_cells.add(current_position)

        if current_position in SPECIAL_CELLS_YELLOW:
            reached_yellow_cell = True
            break

        for action, movement in ACTIONS.items():
            next_row, next_col = current_position[0] + movement[0], current_position[1] +
movement[1]
            next_position = (next_row, next_col)

            if (0 <= next_row < ROWS and 0 <= next_col < COLS and
CELL_VALUES[next_position] != -1 and next_position not in visited_cells):
                queue.append(next_position)
                parent[next_position] = current_position

    # Trace back the path if the agent reached a yellow cell
    if reached_yellow_cell:

```

```

path = []
while current_position:
    path.append(current_position)
    current_position = parent[current_position]

path.reverse() # Reverse the path to start from the agent's position

for next_position in path[1:]: # Exclude the agent's initial position
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            pygame.quit()
            sys.exit()

    # Fill the background
    win.fill(BACKGROUND_COLOR)

    # Draw the grid
    draw_grid()

    # Draw cells
    for row in range(ROWS):
        for col in range(COLS):
            draw_cell(row, col)

    # Draw the agent
    draw_agent(next_position)


    pygame.display.update()

    # Add a delay to visualize movement
    time.sleep(0.5)

# Keep the window open until the user closes it manually
while True:
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            pygame.quit()
            sys.exit()

if __name__ == "__main__":
    main()

```



2.

```
import pygame
import sys
import numpy as np
import time
from collections import deque

pygame.init()

WIDTH, HEIGHT = 500, 500
ROWS, COLS = 10, 10
CELL_SIZE = WIDTH // COLS # Adjusted cell size
GRID_COLOR = (0, 0, 0)
BACKGROUND_COLOR = (255, 255, 255)
SPECIAL_CELL_COLOR_YELLOW = (255, 255, 0)
SPECIAL_CELL_COLOR_RED = (255, 0, 0)
SPECIAL_CELL_COLOR_GREEN = (0, 255, 0)
SPECIAL_CELLS_YELLOW = [(0, 8), (0, 9), (1, 8), (1, 9)]
SPECIAL_CELLS_RED = [(2,2),(2,3),(2,4),(2,5),(2,6),(3,2),(4,2),(5,2),(6,2),(7,2),
                      (8,2),(3,6),(4,6),(5,6),(6,6),(7,6),(8,6),(8,5),(8,3)]
CELL_VALUES = {}

# Set values for special cells
for cell in SPECIAL_CELLS_YELLOW:
    CELL_VALUES[cell] = 20 # Reward for reaching the goal
```

```

for cell in SPECIAL_CELLS_RED:
    CELL_VALUES[cell] = -1 # Penalty for hitting red cells

# Set value for all remaining cells
for row in range(ROWS):
    for col in range(COLS):
        if (row, col) not in CELL_VALUES:
            CELL_VALUES[(row, col)] = 1 # Default value for white cells

# Map actions to movements
ACTIONS = {"up": (-1, 0), "down": (1, 0), "left": (0, -1), "right": (0, 1)}

# Set up the display
win = pygame.display.set_mode((WIDTH, HEIGHT))
pygame.display.set_caption("Reinforcement Learning on Grid")

class Bandit:
    def __init__(self, k, epsilon):
        self.k = k # Number of arms/actions
        self.epsilon = epsilon # Exploration rate
        self.q_estimates = np.zeros(k) # Estimated values of actions
        self.action_counts = np.zeros(k) # Number of times each action is chosen

    def choose_action(self):
        if np.random.rand() < self.epsilon:
            return np.random.choice(range(self.k)) # Choose a random action index with
probability epsilon
        else:
            # Choose the action with the highest estimated value (breaking ties randomly)
            max_actions = np.where(self.q_estimates == np.max(self.q_estimates))[0]
            return np.random.choice(max_actions)

    def update_estimates(self, action_index, reward):
        self.action_counts[action_index] += 1
        # Update the estimated value of the chosen action using the incremental sample
average formula
        self.q_estimates[action_index] += (1 / self.action_counts[action_index]) * (reward -
self.q_estimates[action_index])

    def bandit_environment(self, action_index):
        movement = ACTIONS[list(ACTIONS.keys())[action_index]]
        next_row = 0 + movement[0]
        next_col = 0 + movement[1]

        # Ensure the next position falls within the grid boundaries
        if 0 <= next_row < ROWS and 0 <= next_col < COLS:
            next_position = (next_row, next_col)
            return np.random.normal(loc=CELL_VALUES[next_position], scale=1.0)
        else:

```

```

        # If the next position is out of bounds, return a default reward
        return 0 # You can adjust the default reward as needed

def draw_grid():
    for x in range(0, WIDTH, CELL_SIZE):
        pygame.draw.line(win, GRID_COLOR, (x, 0), (x, HEIGHT))
    for y in range(0, HEIGHT, CELL_SIZE):
        pygame.draw.line(win, GRID_COLOR, (0, y), (WIDTH, y))

def draw_cell(row, col):
    cell_rect = pygame.Rect(col * CELL_SIZE, row * CELL_SIZE, CELL_SIZE, CELL_SIZE)
    if (row, col) in SPECIAL_CELLS_YELLOW:
        pygame.draw.rect(win, SPECIAL_CELL_COLOR_YELLOW, cell_rect)
    elif (row, col) in SPECIAL_CELLS_RED:
        pygame.draw.rect(win, SPECIAL_CELL_COLOR_RED, cell_rect)
    else:
        pygame.draw.rect(win, BACKGROUND_COLOR, cell_rect)

def draw_agent(agent_position):
    agent_rect = pygame.Rect(agent_position[1] * CELL_SIZE, agent_position[0] * CELL_SIZE,
CELL_SIZE, CELL_SIZE)
    pygame.draw.rect(win, SPECIAL_CELL_COLOR_GREEN, agent_rect)

def main(num_steps, k, epsilon):
    agent_position = (4, 4) # Initial position of the agent
    visited_cells = set() # Keep track of visited cells
    queue = deque([agent_position]) # Queue for breadth-first search
    parent = {agent_position: None} # Parent dictionary for storing the shortest path
    reached_yellow_cell = False # Flag to indicate if the agent reached a yellow cell

    bandit = Bandit(k, epsilon)
    rewards = []

    while queue:
        current_position = queue.popleft()
        visited_cells.add(current_position)

        if current_position in SPECIAL_CELLS_YELLOW:
            reached_yellow_cell = True
            break

        for action_index in range(k): # Iterate over all possible action indices
            movement = ACTIONS[list(ACTIONS.keys())[action_index]]
            next_row, next_col = current_position[0] + movement[0], current_position[1] +
movement[1]
            next_position = (next_row, next_col)

```

```

        if (0 <= next_row < ROWS and 0 <= next_col < COLS and
            CELL_VALUES[next_position] != -1 and next_position not in visited_cells):
            queue.append(next_position)
            parent[next_position] = current_position

    # Choose action from Bandit algorithm
    action_index = bandit.choose_action()
    reward = bandit.bandit_environment(action_index)
    bandit.update_estimates(action_index, reward)
    rewards.append(reward)

# Trace back the path if the agent reached a yellow cell
if reached_yellow_cell:
    path = []
    while current_position:
        path.append(current_position)
        current_position = parent[current_position]

    path.reverse() # Reverse the path to start from the agent's position

for next_position in path[1:]: # Exclude the agent's initial position
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            pygame.quit()
            sys.exit()

    # Fill the background
    win.fill(BACKGROUND_COLOR)

    # Draw the grid
    draw_grid()

    # Draw cells
    for row in range(ROWS):
        for col in range(COLS):
            draw_cell(row, col)

    # Draw the agent
    draw_agent(next_position)

    pygame.display.update()

    # Add a delay to visualize movement
    time.sleep(0.5)

# Keep the window open until the user closes it manually
while True:
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            pygame.quit()

```

```
sys.exit()
```

```
if __name__ == "__main__":  
    main(num_steps=100, k=len(ACTIONS), epsilon=0.1)
```

Reinforcement Learning on Grid

